

Biomedical Callsifier

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
import librosa
import os

DATASET_PATH = r"C:\Users\suman\Downloads\ai suman\Dataset Part 1\
Heart_sount_dataset"

data = []
labels = []
sr = 22050

for dataset_name in os.listdir(DATASET_PATH):
    dataset_dir = os.path.join(DATASET_PATH, dataset_name)
    if os.path.isdir(dataset_dir):
        for class_name in os.listdir(dataset_dir):
            class_dir = os.path.join(dataset_dir, class_name)
            if os.path.isdir(class_dir):
                for fname in os.listdir(class_dir):
                    if fname.lower().endswith(".wav"):
                        fpath = os.path.join(class_dir, fname)
                        try:
                            signal, file_sr = librosa.load(fpath,
sr=sr)

                            data.append(signal)
                            labels.append(class_name)
                        except Exception as e:
                            print(f"Error loading {fpath}: {e}")

print(f"Total audio samples loaded: {len(data)}")

if len(data) == 0:
    print("No audio data found. Please check dataset path and
structure.")
else:
    max_len = max(len(s) for s in data)
    padded_data = np.array([np.pad(s, (0, max_len - len(s)),
'constant') for s in data])

    df = pd.DataFrame(padded_data)
    df['label'] = labels
    scaler = MinMaxScaler()
    signal_data = df.drop('label', axis=1).values
    normalized_signal_data = scaler.fit_transform(signal_data.T).T
```

```

df_normalized = pd.DataFrame(normalized_signal_data,
columns=df.columns[:-1])
df_normalized['label'] = labels

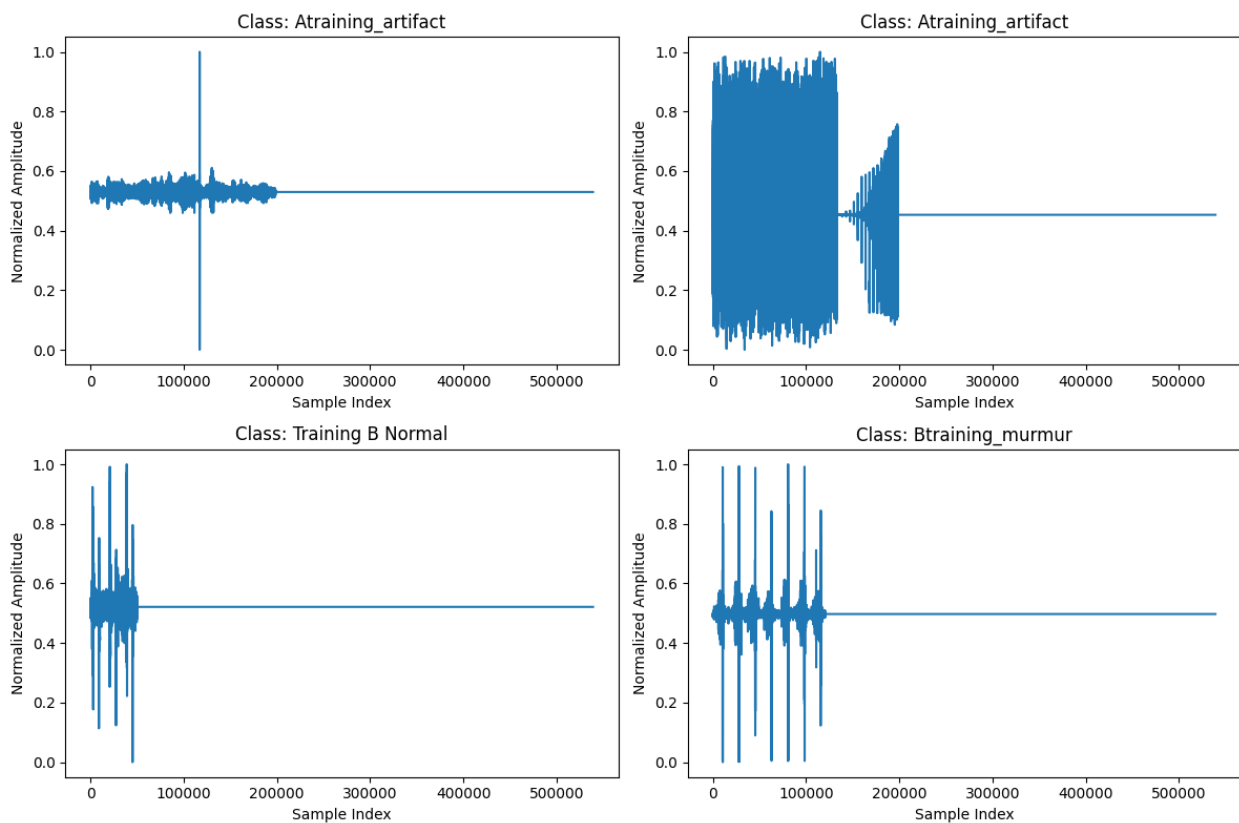
fig, axes = plt.subplots(2, 2, figsize=(12, 8))
axes = axes.flatten()

for i in range(4):
    row_index = np.random.randint(len(df_normalized))
    signal = df_normalized.iloc[row_index, :-1].values
    label = df_normalized.iloc[row_index, -1]
    axes[i].plot(signal)
    axes[i].set_title(f'Class: {label}')
    axes[i].set_xlabel('Sample Index')
    axes[i].set_ylabel('Normalized Amplitude')

plt.tight_layout()
plt.show()
print(df_normalized.head())

```

Total audio samples loaded: 683



6 \ 0 1 2 3 4 5

```

0  0.506175  0.506066  0.506034  0.506041  0.506125  0.506224
0.506226
1  0.521508  0.521879  0.521900  0.521785  0.521716  0.521950
0.521606
2  0.513764  0.508600  0.510990  0.511940  0.514677  0.516584
0.515283
3  0.547782  0.545484  0.546443  0.544580  0.546316  0.550238
0.550870
4  0.494319  0.495040  0.499570  0.503376  0.504970  0.501841
0.500621

      7      8      9  ...  539086  539087  539088
539089  \
0  0.506146  0.506036  0.506028  ...  0.506158  0.506158  0.506158
0.506158
1  0.522069  0.521860  0.521535  ...  0.521842  0.521842  0.521842
0.521842
2  0.512572  0.513687  0.513675  ...  0.512795  0.512795  0.512795
0.512795
3  0.552719  0.553827  0.554481  ...  0.549421  0.549421  0.549421
0.549421
4  0.500215  0.500507  0.500411  ...  0.494241  0.494241  0.494241
0.494241

      539090  539091  539092  539093  539094
label
0  0.506158  0.506158  0.506158  0.506158  0.506158
Atraining_artifact
1  0.521842  0.521842  0.521842  0.521842  0.521842
Atraining_artifact
2  0.512795  0.512795  0.512795  0.512795  0.512795
Atraining_artifact
3  0.549421  0.549421  0.549421  0.549421  0.549421
Atraining_artifact
4  0.494241  0.494241  0.494241  0.494241  0.494241
Atraining_artifact

[5 rows x 539096 columns]

```

Feature extraction

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import librosa
import librosa.display
import soundfile as sf
import os

```

```

import warnings
from IPython.display import display

warnings.filterwarnings("ignore", category=UserWarning)
warnings.filterwarnings("ignore", category=FutureWarning)

dataset_path = r"C:\Users\suman\Downloads\ai suman\Dataset Part 1\
Heart_sount_dataset"

def load_audio(file_path):
    try:
        if not os.path.exists(file_path):
            return None, None
        signal, sr = sf.read(file_path, dtype="float32")
        if len(signal.shape) > 1:
            signal = np.mean(signal, axis=1)
    except:
        try:
            signal, sr = librosa.load(file_path, sr=None, mono=True)
        except:
            return None, None
    return signal, sr

def extract_features(signal, sr=22050):
    mfccs = librosa.feature.mfcc(y=signal, sr=sr, n_mfcc=13,
hop_length=512)
    features = {
        'mfccs_mean': np.mean(mfccs, axis=1),
        'mfccs_std': np.std(mfccs, axis=1),
        'mean': np.mean(signal),
        'std': np.std(signal),
        'max': np.max(signal),
        'min': np.min(signal),
        'peak_amplitude': np.max(np.abs(signal))
    }
    return features

extracted_features_list = []
valid_files = []

for root, dirs, files in os.walk(dataset_path):
    for file in files:
        if file.endswith(".wav"):
            file_path = os.path.join(root, file)
            signal, sr = load_audio(file_path)
            if signal is None or len(signal) == 0:
                continue
            label = os.path.basename(root)
            features = extract_features(signal, sr=sr)
            mfccs_mean_flat = features['mfccs_mean'].flatten()

```

```

        mfccs_std_flat = features['mfccs_std'].flatten()
        feature_row = {f'mfccs_mean_{i}': mfccs_mean_flat[i] for i
in range(len(mfccs_mean_flat))}
        feature_row.update({f'mfccs_std_{i}': mfccs_std_flat[i]
for i in range(len(mfccs_std_flat))})
        feature_row.update({
            'mean': features['mean'],
            'std': features['std'],
            'max': features['max'],
            'min': features['min'],
            'peak_amplitude': features['peak_amplitude'],
            'label': label,
            'file': file
        })
        extracted_features_list.append(feature_row)
        valid_files.append((file_path, label, file))

df_features = pd.DataFrame(extracted_features_list)


if len(valid_files) >= 2:
    fig, axes = plt.subplots(1, 2, figsize=(15, 5))
    for i in range(2):
        file_path, label, file_name = valid_files[i]
        signal, sr = load_audio(file_path)
        if signal is None or len(signal) == 0:
            continue
        mfccs = librosa.feature.mfcc(y=signal, sr=sr, n_mfcc=13,
hop_length=512)
        if mfccs.shape[1] == 0:
            continue
        img = librosa.display.specshow(mfccs, sr=sr, x_axis='time',
ax=axes[i])
        axes[i].set_title(f'MFCCs for {label} ({file_name})')
        fig.colorbar(img, ax=axes[i], format='%+2.0f dB')
        plt.tight_layout()
        plt.show()


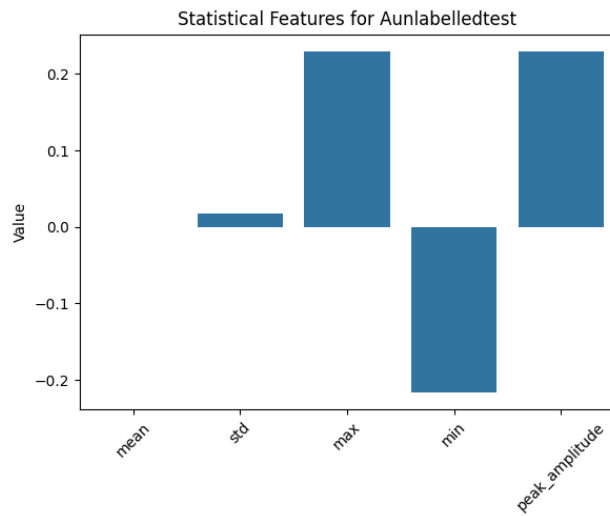
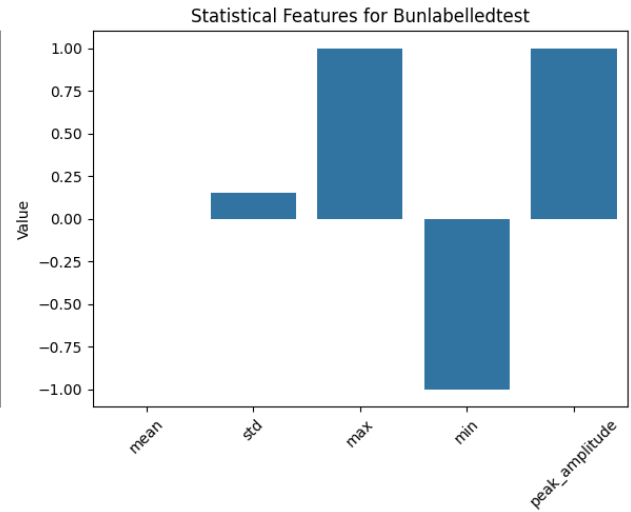
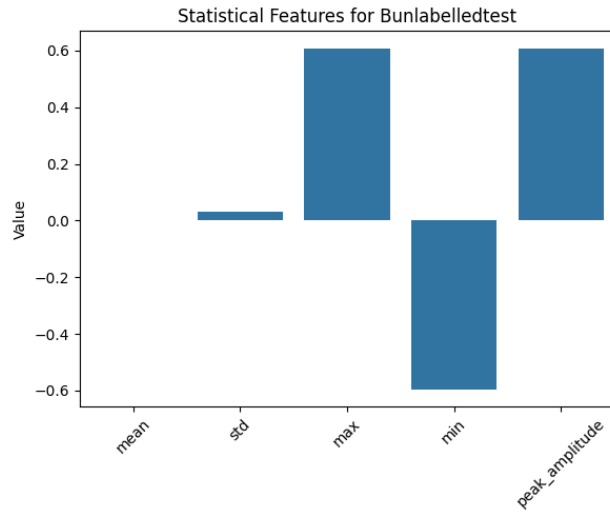
if len(df_features) > 0:
    fig, axes = plt.subplots(2, 2, figsize=(12, 10))
    axes = axes.flatten()
    statistical_features = ['mean', 'std', 'max', 'min',
'peak_amplitude']
    for i in range(4):
        row_index = np.random.randint(len(df_features))
        feature_values = df_features.iloc[row_index]
[statistical_features].values
        label = df_features.iloc[row_index]['label']
        sns.barplot(x=statistical_features, y=feature_values,
ax=axes[i])
        axes[i].set_title(f'Statistical Features for {label}')

```

```
axes[i].set_ylabel('Value')
axes[i].tick_params(axis='x', rotation=45)
plt.tight_layout()
plt.show()
```

```
display(df_features.head())
```





```
{
  "columns": [
    {
      "name": "index",
      "rawType": "int64",
      "type": "integer"
    },
    {
      "name": "mfccs_mean_0",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_1",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_2",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_3",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_4",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_5",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_6",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_7",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_8",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_9",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_10",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_11",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_mean_12",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_std_0",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_std_1",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_std_2",
      "rawType": "float32",
      "type": "float"
    },
    {
      "name": "mfccs_std_3",
      "rawType": "float32",
      "type": "float"
    }
  ]
}
```

```
{ "name": "mfccs_std_4", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_5", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_6", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_7", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_8", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_9", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_10", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_11", "rawType": "float32", "type": "float"},
{ "name": "mfccs_std_12", "rawType": "float32", "type": "float"},
{ "name": "mean", "rawType": "float32", "type": "float"},
{ "name": "std", "rawType": "float32", "type": "float"},
{ "name": "max", "rawType": "float32", "type": "float"},
{ "name": "min", "rawType": "float32", "type": "float"},
{ "name": "peak_amplitude", "rawType": "float32", "type": "float"},
{ "name": "label", "rawType": "object", "type": "string"},
{ "name": "file", "rawType": "object", "type": "string"}], "ref": "a4601573-6251-48dc-b8b2-e28e2b8238db", "rows": [{"0", "-599.93896", "33.781063", "7.35434", "4.5494447", "2.182476", "2.7137213", "0.70974946", "-0.5813915", "-1.2105308", "-2.7702992", "-1.950883", "-1.373695", "-1.6228255", "63.559986", "40.340088", "13.842911", "12.959423", "9.30077", "7.30098", "5.678077", "5.4349117", "4.6704197", "4.63749", "3.8292885", "3.345399", "3.240959", "-1.0466322e-05", "0.014274816", "0.70410156", "-0.7289734", "0.7289734", "Atraining_artifact", "201012172012.wav"}, {"1", "-752.226", "23.831118", "8.511935", "-12.289284", "6.590361", "-13.384438", "8.085058", "-10.728707", "8.7330475", "-9.553118", "3.778762", "-4.386584", "-2.2409499", "193.43013", "67.308914", "27.99302", "15.005083", "11.636027", "4.58469", "5.4975467", "4.2518296", "5.7577662", "2.9655578", "3.7120292", "4.586085", "3.1971843", "-1.1055222e-06", "0.00084005133", "0.016021729", "-0.018188477", "0.018188477", "Atraining_artifact", "201105040918.wav"}, {"2", "-587.1053", "56.313488", "-21.25168", "5.092139", "-13.107715", "-3.1940339", "-3.2685952", "-9.355012", "3.2398045", "-8.664276", "1.9791583", "-5.9559555", "-1.5866339", "78.01337", "21.039032", "9.313816", "12.808322", "11.165461", "10.16053", "6.3035564", "4.477588", "5.2280607", "4.9290957", "4.7142653", "4.2985764", "4.010035", "4.744103e-07", "0.0010501887", "0.020721436", "-0.027069092", "0.027069092", "Atraining_artifact", "201105041959.wav"}, {"3", "-349.67032", "141.82634", "-21.67626", "23.579344", "-2.894292", "-3.9263988", "-8.095814", "-8.885119", "-7.6928396", "-5.4134545", "-5.656429", "-9.26078", "-6.8140674", "65.990234", "31.441687", "24.155802", "12.032792", "11.656161", "9.121105", "7.678339", "7.4909124", "6.634921", "6.0082426", "6.1362395", "6.3897085", "5.5943704", "-1.3695388e-05", "0.017936688", "0.88171387", "-0.90515137", "0.90515137", "Atraining_artifact", "201105051017.wav"}, {"4", "-333.93628", "115.90495", "-7.5759363", "5.6850233", "1.2802495", "-1.1377603", "2.2690852", "-7.8378334", "-2.7766035", "-10.6930485", "-0.80935055", "-7.148301", "-
```



```
5.4792757", "85.54612", "22.55842", "16.148848", "8.672056", "9.865658", "7.1070833", "6.7316422", "7.4879208", "6.4449434", "6.3067145", "6.4818907", "6.2358828", "7.083377", "-2.8633624e-05", "0.038326383", "0.886261", "-0.94033813", "0.94033813", "Atraining_artifact", "201105060108.wav"]], "shape": {"columns": 33, "rows": 5}}
```

Data splitting

```
from sklearn.model_selection import train_test_split

X = df_features.drop('label', axis=1)
y = df_features['label']
X_train, X_temp, y_train, y_temp = train_test_split(X, y,
test_size=0.3, random_state=42, stratify=y)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp,
test_size=0.5, random_state=42, stratify=y_temp)

print("Training set shape (X_train, y_train):", X_train.shape,
y_train.shape)
print("Validation set shape (X_val, y_val):", X_val.shape,
y_val.shape)
print("Test set shape (X_test, y_test):", X_test.shape, y_test.shape)

Training set shape (X_train, y_train): (582, 32) (582,)
Validation set shape (X_val, y_val): (125, 32) (125,)
Test set shape (X_test, y_test): (125, 32) (125,)
```

Model building

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

logistic_regression_model = LogisticRegression(random_state=42)
random_forest_model = RandomForestClassifier(random_state=42)
svm_model = SVC(random_state=42)

dnn_model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    Dense(32, activation='relu'),
    Dense(1, activation='sigmoid')
])

dnn_model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
```

```

print("Logistic Regression Model:")
print(logistic_regression_model)
print("\nRandom Forest Model:")
print(random_forest_model)
print("\nSVM Model:")
print(svm_model)
print("\nDeep Neural Network Model Summary:")
dnn_model.summary()

```

Logistic Regression Model:
LogisticRegression(random_state=42)

Random Forest Model:
RandomForestClassifier(random_state=42)

SVM Model:
SVC(random_state=42)

Deep Neural Network Model Summary:

Model: "sequential"

Layer (type) Param #	Output Shape	
dense (Dense) 2,112	(None, 64)	
dense_1 (Dense) 2,080	(None, 32)	
dense_2 (Dense) 33	(None, 1)	

Total params: 4,225 (16.50 KB)

Trainable params: 4,225 (16.50 KB)

Non-trainable params: 0 (0.00 B)

Model training and evaluation

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

```

```

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE

sm = SMOTE(random_state=42)
X_train_res, y_train_res = sm.fit_resample(X_train, y_train)

models = {
    "Logistic Regression": LogisticRegression(max_iter=2000,
class_weight="balanced"),
    "Random Forest": RandomForestClassifier(n_estimators=300,
max_depth=20, class_weight="balanced"),
    "SVM (RBF)": SVC(kernel="rbf", C=10, gamma='scale',
probability=True, class_weight="balanced")
}

for name, model in models.items():
    model.fit(X_train_res, y_train_res)
    y_pred = model.predict(X_test)

    print(f"\n{name} Results")
    print("Accuracy:", round(accuracy_score(y_test, y_pred), 4))
    print("\nClassification Report:\n", classification_report(y_test,
y_pred, zero_division=0))


    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(5,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f"{name} Confusion Matrix")
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.show()

```

Logistic Regression Results
Accuracy: 0.6168

Classification Report:

	precision	recall	f1-score	support
0	0.53	0.69	0.60	70
1	0.71	0.57	0.63	97
accuracy			0.62	167
macro avg	0.62	0.63	0.62	167
weighted avg	0.64	0.62	0.62	167




Random Forest Results

Accuracy: 0.6707

Classification Report:

	precision	recall	f1-score	support
0	0.60	0.63	0.62	70
1	0.72	0.70	0.71	97
accuracy			0.67	167
macro avg	0.66	0.66	0.66	167
weighted avg	0.67	0.67	0.67	167




SVM (RBF) Results

Accuracy: 0.5569

Classification Report:

	precision	recall	f1-score	support
0	0.47	0.53	0.50	70
1	0.63	0.58	0.60	97
accuracy			0.56	167
macro avg	0.55	0.55	0.55	167
weighted avg	0.56	0.56	0.56	167



Model Validation

```
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report, confusion_matrix,
roc_auc_score
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
logistic_regression_model = LogisticRegression(max_iter=2000,
class_weight="balanced")
logistic_regression_model.fit(X_train, y_train)

random_forest_model = RandomForestClassifier(n_estimators=300,
max_depth=20, class_weight="balanced")
random_forest_model.fit(X_train, y_train)

svm_model = SVC(kernel="rbf", C=10, gamma='scale', probability=True,
class_weight="balanced")
svm_model.fit(X_train, y_train)

models = {
    "Logistic Regression": logistic_regression_model,
    "Random Forest": random_forest_model,
    "SVM": svm_model
}
```

```

for name, model in models.items():
    print(f"\n {name} Evaluation")

    y_pred = model.predict(X_test)
    print("Accuracy:", round(accuracy_score(y_test, y_pred), 4))
    print("Precision:", round(precision_score(y_test, y_pred,
average='weighted', zero_division=0), 4))
    print("Recall:", round(recall_score(y_test, y_pred,
average='weighted', zero_division=0), 4))
    print("F1-score:", round(f1_score(y_test, y_pred,
average='weighted', zero_division=0), 4))
    print("\nClassification Report:\n", classification_report(y_test,
y_pred, zero_division=0))

    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f"{name} Confusion Matrix")
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.show()

    if hasattr(model, "predict_proba"):
        y_score = model.predict_proba(X_test)
        if len(np.unique(y_test)) == 2:
            auc = roc_auc_score(y_test, y_score[:,1])
        else:
            auc = roc_auc_score(pd.get_dummies(y_test), y_score,
multi_class='ovr', average='weighted')
        print("ROC-AUC:", round(auc, 4))

```

Logistic Regression Evaluation

Accuracy: 0.6168


Precision: 0.6384

Recall: 0.6168

F1-score: 0.6187

Classification Report:

	precision	recall	f1-score	support
0	0.53	0.69	0.60	70
1	0.71	0.57	0.63	97
accuracy			0.62	167
macro avg	0.62	0.63	0.62	167
weighted avg	0.64	0.62	0.62	167



ROC-AUC: 0.6736

□ Random Forest Evaluation

Accuracy: 0.6407


Precision: 0.6476

Recall: 0.6407

F1-score: 0.6427

Classification Report:

	precision	recall	f1-score	support
0	0.56	0.63	0.59	70
1	0.71	0.65	0.68	97
accuracy			0.64	167
macro avg	0.64	0.64	0.64	167
weighted avg	0.65	0.64	0.64	167



ROC-AUC: 0.7041

□ SVM Evaluation

Accuracy: 0.5569


Precision: 0.5643

Recall: 0.5569

F1-score: 0.5593

Classification Report:

	precision	recall	f1-score	support
0	0.47	0.53	0.50	70
1	0.63	0.58	0.60	97
accuracy			0.56	167
macro avg	0.55	0.55	0.55	167
weighted avg	0.56	0.56	0.56	167



ROC-AUC: 0.6358

Confusion matrix and visualizations

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
import numpy as np

all_models = {
    'Logistic Regression': logistic_regression_model,
    'Random Forest': random_forest_model,
    'SVM': svm_model
}

fig, axes = plt.subplots(1, len(all_models), figsize=(18, 5))

for ax, (name, model) in zip(axes, all_models.items()):
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax)
    ax.set_title(name)
    ax.set_xlabel('Predicted')
    ax.set_ylabel('True')
```

```
plt.tight_layout()
plt.show()
```

